# Honest Al: Fine-Tuning "Small" Language Models to Say "I Don't Know", and Reducing Hallucination in RAG

Xinxi Chen\* Independent Researcher xc336@cornell.edu Li Wang\* Independent Researcher li@liwang.info

Wei Wu\* Independent Researcher drweiwu@outlook.com

Qi Tang\* Independent Researcher qitang@closeby-ai.com Yiyao Liu\* Independent Researcher info@yiyaoliu.com

# **Abstract**

Hallucination is a key roadblock for applications of Large Language Models (LLMs), particularly for enterprise applications that are sensitive to information accuracy. To address this issue, two general approaches have been explored: Retrieval-Augmented Generation (RAG) to supply LLMs with updated information as context, and fine-tuning the LLMs with new information and desired output styles. In this paper, we propose Honest AI: a novel strategy to fine-tune "small" language models to say "I don't know" to reduce hallucination, along with several alternative RAG approaches. The solution ranked 1st in Task 2 for the false premise question<sup>1</sup>. The alternative approaches include using RAG with search engine and knowledge graph results, fine-tuning base LLMs with new information and combinations of both approaches. Although all approaches improve the performance of the LLMs, RAG alone does not significantly improve the performance and fine-tuning is needed for better results. Finally, the hybrid approach achieved the highest score in the CRAG benchmark [27]. In addition, our approach emphasizes the use of relatively small models with fewer than 10 billion parameters, promoting resource efficiency.

#### **CCS** Concepts

• Computing methodologies → Natural language generation.

# Keywords

Large Language Models (LLM), Retrieval Augmented Generation (RAG), Knowledge Graph, Search

#### **ACM Reference Format:**

Xinxi Chen, Li Wang, Wei Wu, Qi Tang, and Yiyao Liu. 2024. Honest AI: Fine-Tuning "Small" Language Models to Say "I Don't Know", and Reducing Hallucination in RAG. In Proceedings of 2024 KDD Cup Workshop for Retrieval Augmented Generation at the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDDCup '24). ACM, New York, NY, USA, 8 pages.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

KDDCup '24, Aug 25-29, 2024, Barcelona, Spain

© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.

## 1 Introduction

Large Language Models (LLMs), as a type of foundation models with general language capabilities, have eclipsed traditional Natural Language Processing (NLP) models that focus on specific tasks in majority of NLP applications since the inception of GPT [1]. Supervised fine tuning (SFT) using labeled data and reinforcement learning from human feedback (RLHF) using preference data have proven effective in further enhancing LLMs' performance and alignment (e.g., ChatGPT [17] for question answering applications). However, LLMs suffer from hallucination, which hinders their application in accuracy-sensitive scenarios, such as the enterprise applications.

To alleviate the hallucinations of LLMs, several approaches have been proposed, including retrieval-augmented generations (RAGs) [12] and fine-tuning with domain-specific knowledge. In this paper, we summarize the approaches we tried in the 2024 Meta KDD Cup competition with Comprehensive RAG Benchmark (CRAG) data [27]. Since the CRAG benchmark focuses on difficult problems for LLMs, vanilla LLMs fail to perform well right out of the box. It turns out that RAG alone is not enough to alleviate hallucination in the benchmark and fine-tuning is needed to achieve higher accuracy. Our results show that the hybrid approach using both RAG and fine tuning performs best in CRAG. With those approaches combined, our team, Team Future, ranked high in each task of the competition and won first place in the false premise question in Task 2 (Fig. 1).

The review of related work is summarized in section 2, followed by a description of methodologies in section 3. The results are shown in section 4, with conclusions and future works discussed in section 5 and 6.

#### Task 2: Knowledge Graph and Web Retrieval

- Team db3
   Team APEX
- 🥉 Team md\_dh
- A Simple\_w\_condition Question: Team ElectricSheep
- A Set Question: Team ElectricSheep
- A Comparison Question: Team dRAGonRAnGers
- Aggregation Question: Team ElectricSheep
- A Multi-hop Question: Team ElectricSheep
- 🔏 Post-processing Question: Team ElectricSheep
- 🗸 False Premise Question: Team Future

Figure 1: Team Future in Meta CRAG Challenge 2024 Winner List

<sup>\*</sup>These authors contributed equally to this work. Emails are provided for contact purposes only and do not represent an endorsement by any institution.

<sup>&</sup>lt;sup>1</sup>The Meta CRAG Challenge has over 2000 participants with over 5500 submissions.

#### 2 Related Work

LLMs are very good at memorizing the content used for their pretraining. However, there are many drawbacks to using the memorization capability directly for Question Answering (QA) tasks. For example, depending on the size of the model, the quality of the pretraining data, and the type of questions, LLMs' memorization capability can be very limited and difficult to control. LLMs are also challenging to update except through retraining or fine-tuning, so they cannot handle questions on recent events if deployed. Arguably, the most problematic drawback of using LLMs for QA is that they can hallucinate, especially when the models are unsure if the context contains the information needed to answer a question.

Fine-tuning is a straightforward way to update the knowledge of LLMs when new information becomes available. However, due to the scarcity of GPUs and limited access to high-quality data, it is not feasible for many use cases. Furthermore, the behavior of fine-tuning language models is not well studied, and fine-tuning with poor data or practices might decrease the model's capabilities, causing modality collapse unless more complex methods such as RLHF are used [17].

Retrieval-Augmented Generation (RAG) [8, 12] is a popular method to address the shortcomings of LLMs, by augmenting them with non-parametric data sources, and leveraging LLMs' powerful in-context learning [2] capability. [8] grouped the approaches of using RAG for LLMs into three categories: Naive RAG, Advanced RAG [14, 31], and Modular RAG [22, 29]. Naive and Advanced RAG approaches are widely used in practice due their simplicity and low development cost. These approaches generally consist of three parts: curating non-parametric databases, retrieving relevant snippets from the databases given the query, and generating responses using LLMs through in-context learning and prompt engineering with the related snippets.

While the research on combining LLMs and RAG for QA mainly focuses on text, there is also research exploring the use of resources beyond text, such as images [4], audio [30], video [26], and code [16] to enhance the capabilities of language models.

There are various efforts to create RAG benchmarks and proposing appropriate evaluation metrics in recent years. [27] is one of the recent ones, which forms the foundation of this paper. [27] created a factual question answering benchmark of 4,409 question-answer pairs and mock APIs to simulate web and Knowledge Graph (KG) searches. It also proposes an evaluation mechanism that distinguishes between hallucinations and missing answers, and assigns a higher penalty to hallucinations. [3] created a RAG evaluation benchmark in both English and Chinese, and analyzed different LLMs from 4 aspects: noise robustness, negative rejection, information integration, and counterfactual. They found that LLMs demonstrate a certain degree of noise robustness, but struggle significantly in other aspects.

Apart from specific RAG datasets, there are many existing QA datasets that include context passages for each question. These datasets can also be used for RAG experiments, and cover a wide range of questions such as multiple-choice QA [5, 18, 23], single-hop QA [10, 11, 19], multi-hop QA [9, 24, 28], and domain-specific QA [6, 15, 25].

# 3 Methodologies

In this section, we will first introduce the CRAG dataset and then describe the main approaches we have tried, including native RAG, fine-tuning, and hybrid approaches.

#### 3.1 Dataset

The CRAG dataset includes five question domains (e.g. movie, sports, and etc.) with varying levels of complexity, ranging from straightforward facts to those requiring reasoning (e.g. false premise and multi-hop reasoning) [27]. It also considers facts with different levels of timeliness (e.g. real-time, fast-changing, slow-changing, and stable). Figure 2 illustrates the distribution of question types across different domains and timeliness categories. Notably, the characteristics of question distributions in the movie and finance domains differ significantly. Real-time and fast-changing questions necessitate access to relevant, up-to-date data sources and effective RAG implementations. In contrast, the movie domain contains more static questions than the finance domain, which may be easier to address.

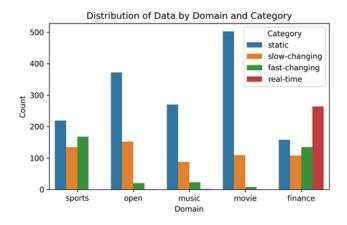


Figure 2: The Distribution of Questions in Different Domains and Types of Timeliness

Our solution ranked first in Task 2 for false premise questions. A false premise question is defined as a question with a false assumption. For example, "What's the name of Taylor Swift's rap album before she transitioned to pop?" This is a false premise question because Taylor Swift didn't release any rap albums, and the expected answer is "invalid question" [27].

In this competition, we chose a divide-and-conquer approach to design a RAG system to tackle different types of questions. We assumed that categorizing these questions could be done easily in practice. The objective is to identify the key parametric differences, such as varying prompts, parsing techniques, chunking sizes, top-k thresholds, and rule designs, to optimize performance and achieve a higher score during implementation.

#### 3.2 Scoring

The overall score is a macro-average across all domains. The scoring system awards scores based on the quality of the response, and penalizes hallucination [27]. If the response is perfect, it receives

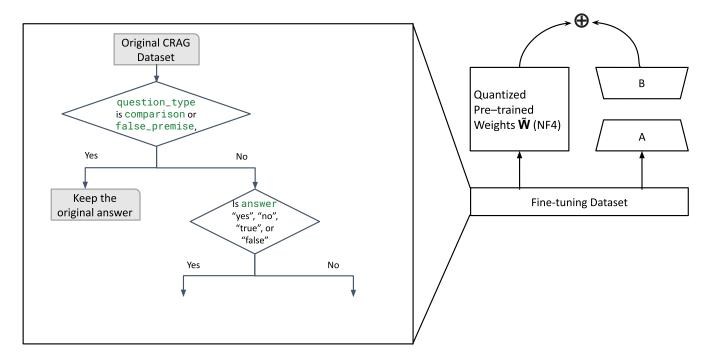


Figure 3: Fine-tuning Process

1 point. If the response is acceptable, it receives 0.5 points. If the response is missing, i.e. "I don't know", it receives 0 points. If the response is incorrect, it receives -1 point.

# 3.3 RAG

RAG [8, 12] is a popular approach to alleviate the hallucinations of LLMs. There are various architectures for RAG in real applications, including advanced RAG techniques. In this paper, we explored several naive RAG approaches. Generally, naive RAG selects the highest cosine similarity results from a vector database and supplies the context as inputs for LLMs. There are multiple ways to find the relevant information to implement RAG. We tried using naive cosine similarity and using LLMs as retrieval and ranking models to find relevant information for each web page.

Furthermore, as an extreme case, we also tried using the state-of-the-art Gemini 1.5 pro model with a 1 million token context window [20] to supply all the retrieved web pages to LLMs. These models with long context windows are promising because they eliminate the need to truncate information from retrieved results. However, with vanilla Gemini 1.5 pro and raw retrieved web pages, we achieved similar results like vanilla RAG with severe hallucination, which is disappointing. Therefore, we didn't further investigate more advanced RAG techniques. An example of hallucination is shown in Appendix A.2.

These results for RAG are summarized in Section 4.

# 3.4 Fine-tuning

Our initial investigation revealed that answers in most categories are very challenging, and the evaluation metric heavily penalizes hallucinations. Therefore, it is better for the model to be "honest" about its limits by replying "i don't know" in cases of uncertainty, rather than providing wrong answers. However, it is non-trivial to assess LLMs' confidence level reliably [13]. Instead, we hypothesize that by explicitly teaching LLMs to reply "i don't know" to challenging questions while providing real answers for easy questions, LLMs may be able to learn the ability to distinguish between challenging and easy questions.

To test this hypothesis, we decided to use the QLoRA [7] technique to fine-tune the Llama-2-7b-chat in 4-bit precision and optimize VRAM usage, due to limited GPU resources. More specifically, as it is shown in Fig. 3, we used the training data provided by the organizer, reserved 250 instances for testing, and made a few modifications to the rest of the data. If the question\_type is comparison or false\_premise, or the answer is "yes", "no", "true" or "false", we do not modify the answers to the questions; otherwise, we replace the original answer with "i don't know". Then we used the modified data to fine-tune the model.

We use Alpha=16, r=64, and Dropout=0.1 for QLoRA. Additionally, we use a batch size of 8, a learning rate of 0.0002, a weight decay of 0.001, and fine-tune the model for 5 epochs. We evaluate the fine-tuned model on the 250 withheld questions, using the offline evaluation script provided by the organizer.

We also experimented with Meta-Llama-3-8B-Instruct, which performed consistently worse.

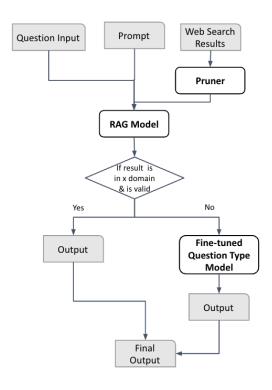


Figure 4: Hybrid Approach

# 3.5 Hybrid Approach

With the fine-tuned model mentioned in the above section, we achieved high rankings in the competition, and first place for the false premise question type in Task 2: Knowledge Graph and Web Retrieval. However, the method did not fully utilize the additional knowledge from the web search results and the knowledge graph. To improve the results, we developed the following hybrid approach to better utilize the knowledge in certain domains.

The hybrid approach leveraged the benefits of both the RAG model and the Fine-tuned Question Type model (Fig. 4). The hybrid approach first utilized the vanilla Meta-Llama-3-8B-Instruct to serve as the RAG model to generate results that include domain and answer information. The next step is to determine if the domain belongs to the movie domain, which has a lower level of hallucination based on the CRAG benchmark results [27]. If it is in the movie domain and the answer is valid, the answer is used as the final answer. If the answer is invalid, the question is sent to the Fine-tuned Question Type Model, which is good at answering false premise questions, and more modest by answering "i don't know" for other types of questions. Specifically, an answer is considered "invalid" in two scenarios: when the response is "invalid question," or when there is a JSON processing error.

For the input of the RAG model, we also applied a Pruner to extract the top k sentences from the web search results. This pruner computes the cosine similarity of each sentence  $W_{ij}$  in the document with the question  $Q_i$ , after converting them to Sentence-BERT embeddings [21]. For each top sentence, the following  $\mathfrak m$  sentences

in the paragraph are appended to the context to enrich the information. To ensure content quality, an additional cosine similarity threshold n is applied, and the answer is used only if this threshold is met for top k sentences.

$$\label{eq:cosine_similarity} \begin{aligned} & \text{cosine\_similarity}(Q_i, W_{ij}) = \frac{Q_i \cdot W_{ij}}{\|Q_i\| \|W_{ij}\|} \end{aligned} \tag{1}$$

Based on offline evaluation, this approach improved the total score from 0.073 to 0.86 using results from 300 samples, compared with the Fine-tuned Question Type model. Since our Fine-tuned Question Type model achieved a score of 0.0960 in the online evaluation for Task 3 (without the holdout test), this hybrid approach is expected to achieve a higher score. It has not been evaluated online yet, because the online evaluation system was closed after phase 2.

In addition, the prompt in Appendix A.1 is used to get the domain name of the question in a JSON output, which can be used for easy processing.

# 4 Results

We firstly tested the Llama3 8b pretrained model with or without RAG and prompt tuning in the domain of movie vs. finance and types of questions of simple vs. post-processing and multi-hop. The results of which are shown in Table 1. Each test case ran 100 samples out of the 2.7k samples in CRAG. Surprisingly, we observed that adding more detailed instructions in the prompt actually dropped the overall performance significantly on the Llama3 8b pretrained model. One hypothesis is that the model's performance is highly sensitive to the format of prompting and needs to be properly configured and fine-tuned.

Given that the retrieval might not be correct, we also tried Gemini 1.5 pro with a 1 million token context window as a long context window models to see if feeding all the retrieved information to the LLMs would perform better than any RAG approaches. The results show no improvement, and we didn't further investigate on this.

Table 2 shows the overall results from our fine-tuned model, which achieved 0.096 with 323 samples from the online judging system. With this model, we achieved the highest score in Task 2 for the false premise problems (Fig. 5). Finally, our hybrid approach results with cosine similarity threshold of 0.75 show that the score improved by 0.013 from 0.073 to 0.086 (Table 3) with fine-tuning and RAG combined. The accuracy increased by 0.026 and there was a slight increase in hallucination by 0.013.

Furthermore, 4 shows some key examples of differences in predictions comparing the fine-tuned only model and the hybrid approach. For the first false premise question "when did hamburg become the biggest city of germany", both models provide the answer "invalid question", since the largest city in Germany is Berlin. For the second question which is a comparison question, both models also provide the same correct answer. For the third question related to movie domain, the fine-tuned model responds with "i don't know", while the hybrid approach provides "inception" which is a correct answer. For the forth question related to movie domain, the fine-tuned model provides "i don't know again", while the hybrid approach provide a relevant answer including who are professionals relevant to the

<sup>&</sup>lt;sup>1</sup>The prompt tuning strategies detailed in Table 1 predominantly follow instructions such as: answer the question given the context, and regulate the answer format. Prompt examples: answer "i don't know" directly, if ..., or "invalid question", if ...

Domain or Question Type	Prompt Tuning	RAG	Accuracy	Hallucination	Score
Movie	n/a	w/o	0.54	0.46	0.08
Movie	n/a	$\mathbf{w}/$	0.52	0.48	-0.04
Movie	Y*1	w/o	0.26	0.41	-0.32
Movie	Y	$\mathbf{w}/$	0.35	0.55	-0.19
Finance	n/a	w/o	0.31	0.69	-0.38
Finance	n/a	$\mathbf{w}/$	0.34	0.66	-0.32
Finance	Y	w/o	0.1	0.41	-0.31
Finance	Y	$\mathbf{w}/$	0.19	0.55	-0.36
Simple	Y	w/o	0.15	0.39	-0.24
Simple	Y	$\mathbf{w}/$	0.24	0.5	-0.26
Post-processing and multi-hop	Y	w/o	0.12	0.5	-0.38
Post-processing and multi-hop	Y	$\mathbf{w}/$	0.24	0.59	-0.35

Table 1: Preliminary test results of Q&A in specific categories

movie (e.g. the director Steve Carr), but not the complete list of the producers.

Task	<b>Question Type</b>	Team	Score
Task 2	simple_w_condition	ElectricSheep	23.9
	set	ElectricSheep	36.65
	comparison	dRAGonRAnGers	38
	aggregation	ElectricSheep	18.75
	multi_hop	ElectricSheep	23.2
	post_processing	ElectricSheep	11.75
	false_premise	Future	64.6

Figure 5: Team Future Got 64.6% Score in false\_premise Question Type

#### 5 Discussion

We tried various approaches to alleviate the hallucinations of LLMs, including Retrieval-Augmented Generations (RAGs) and fine-tuning with domain-specific knowledge, while participating in the 2024 Meta KDD Cup with Comprehensive RAG Benchmark (CRAG) data. It turns out that RAG alone is not enough to perform well in the benchmark, and fine-tuning is needed to achieve a higher score. Our results show that the hybrid approach using both RAG and fine-tuning performs best in CRAG.

For the basic RAG approach, we are uncertain whether the model's accurate answers stem from the pretrained model's prior knowledge or the retrieved data from the reference. At first glance, it appears that basic RAG does not significantly improve the final score. One possible reason is that the retrieved content, while relevant (cosine similarity score > 0.7), might not be useful because it is too basic, general, or vague. Additionally, improper prompting significantly reduces accuracy compared to using no prompts, likely because the pretrained model is misled by the instruction

and loses its own prior knowledge to answer the question. More advanced RAG and prompt design should be tested to draw definitive conclusions.

With further improvement in the hybrid RAG approach, focusing on the movie domain, it is clear that better quality of retrieved content helps improve the RAG results. In the Table 5, the second question shows that if the birth date is correct in the retrieved content, the prediction of Hybrid RAG Approach is correct. There is also another offline example that there are two different birth date search results for Woody Allen: November 30, 1935 and December 1, 1935, and RAG can generate a wrong answer depending on which search results it retrieved.

The CRAG question dataset also contains many questions like "In year x, which type y movie was recognized with the best type y movie in Oscar". a common hallucination type occurs in this way: it provides a movie which was released in year x, but it won Oscar in year x+1. In the table 5, it shows an example that "the incredibles" was provided by RAG as the best animated feature film in Oscar 2004. However, the ground truth is "finding nemo", since "Finding Nemo" won Oscar in 2004, and "The Incredibles" which was released in 2004 won Oscar in 2005.

#### 6 Future Work

With hallucination being one of the biggest challenges in applying LLMs to real-world applications, addressing different types of hallucination continues to require deeper and more extensive research and efforts. Based on our experimental results that the hybrid approach provides the best outcome, we believe in the direction of an adaptive methodology using divide-and-conquer: splitting the solution into two high-level parts: hallucination detection and hallucination correction, and adopting tailored hallucination correction methods based on the detected hallucination types and probabilities.

In terms of detection, questions with low probabilities of hallucination can be questions in common domains, requiring simple logic, and/or about a past fact that doesn't change; questions with high probabilities of hallucination can often be in specialized domains, requiring complex reasoning and multi-hop steps, and/or

Table 2: Fine-tuned Only Model with 323 Samples from Online Judging System (Task 3)

	Exact Accuracy	Accuracy	Hallucination	Missing	Total Score
Fine-tuned Only	0.111	0.152	0.056	0.793	0.096

Table 3: Score Comparison of Fine-tuned Only Model vs. Hybrid Approach with 300 Offline Samples (Task 1)

	Exact Accuracy	Accuracy	Hallucination	Missing	Total Score
Fine-tuned Only	0.107	0.11	0.037	0.853	0.073
Hybrid Approach (cosine similarity threshold = 0.8)	0.117	0.123	0.043	0.833	0.080
Hybrid Approach (cosine similarity threshold = 0.75)	0.113	0.136	0.05	0.813	0.086

Table 4: Answer Comparison of Fine-tuned Only Model vs. Hybrid Approach with 300 Offline Samples (Task 1)

Question	Ground Truth	Fine-tuned Only Model	Hybrid Approach
when did hamburg become the biggest city of germany?	invalid question	invalid question	invalid question
which wta player had a higher singles ranking to end last year, madison keys or daria kasatkina?	madison keys	madison keys	madison keys
what 2010 film was directed by christopher nolan?	inception	i don't know	inception
who were the producers of the movie paul blart: mall cop?	adam sandler, jack giarraputo, kevin james, todd garner, barry bernardi	i don't know	steve carr, kevin james, and nick bakay

Table 5: RAG Results are Sensitive to the Quality of Retrieved Content

Question	Ground Truth	Fine-tuned Only Model	Hybrid RAG Approach	Relevant Retrieved Content
in 2004, which animated film was recognized with the best animated feature film oscar?	finding nemo	i don't know	the incredibles	<doc> Superhero Best Animated Movie Oscar The Incredibles (2004) </doc>
when was the birth of michael bay?	1965-02-17	i don't know	michael bay was born on february 17, 1965.	<doc> What is the Michael Bay date of birth? The DOB for Michael Bay was 17 Feb 1965 </doc>

fast-changing facts. Notice that the hallucination types and probabilities are a full spectrum, and they don't simply depend on the question itself, but also on the external context. For example, the question "which country wins the most gold medals in Olympics Games Paris" is fast-changing while the game is ongoing and will become a past fact after the game is over.

In terms of correction, different strategies of answer correction can be adopted based on the detection results. For example, a realtime question will require real-time querying of the world knowledge (e.g. search engines that index the web in real-time); a specialized domain question could leverage external knowledge.

While both detection and correction remain extremely challenging, the proposed adaptive approaches using divide-and-conquer remain a promising direction for solving the hallucination problems

in LLMs. In addition, due to limited time during the competition, the movie was primarily explored in the RAG hybrid approach, and other domains could also benefit from this approach.

# 7 Acknowledgments

Great thanks to Ermo Wei, who shared invaluable LLM fine-tuning tips that significantly aided our fine-tuning efforts.

### References

- [1] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. Advances in neural information processing systems 33 (2020), 1877–1901.
- [2] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda

- Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems*, Vol. 33. Curran Associates, Inc., 1877–1901.
- [3] Jiawei Chen, Hongyu Lin, Xianpei Han, and Le Sun. 2024. Benchmarking Large Language Models in Retrieval-Augmented Generation. Proceedings of the AAAI Conference on Artificial Intelligence 38, 16 (Mar. 2024), 17754–17762.
- [4] Wenhu Chen, Hexiang Hu, Xi Chen, Pat Verga, and William Cohen. 2022. MuRAG: Multimodal Retrieval-Augmented Generator for Open Question Answering over Images and Text. In Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 5558–5570.
- [5] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. Think you have Solved Question Answering? Try ARC, the AI2 Reasoning Challenge. https://arxiv.org/abs/1803.05457
- [6] Pradeep Dasigi, Kyle Lo, Iz Beltagy, Arman Cohan, Noah A. Smith, and Matt Gardner. 2021. A Dataset of Information-Seeking Questions and Answers Anchored in Research Papers. In Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies. Association for Computational Linguistics.
- [7] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. QLoRA: Efficient Finetuning of Quantized LLMs. In Advances in Neural Information Processing Systems, Vol. 36. Curran Associates, Inc., 10088–10115.
- [8] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Ji-awei Sun, Meng Wang, and Haofen Wang. 2024. Retrieval-Augmented Generation for Large Language Models: A Survey. https://arxiv.org/abs/2312.10997
- [9] Xanh Ho, Anh-Khoa Duong Nguyen, Saku Sugawara, and Akiko Aizawa. 2020. Constructing A Multi-hop QA Dataset for Comprehensive Evaluation of Reasoning Steps. In Proceedings of the 28th International Conference on Computational Linguistics. International Committee on Computational Linguistics, Barcelona, Spain (Online).
- [10] Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. 2017. TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension. In Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). Association for Computational Linguistics, Vancouver, Canada, 1601–1611.
- [11] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. 2019. Natural Questions: A Benchmark for Question Answering Research. Transactions of the Association for Computational Linguistics 7 (08 2019), 453–466.
- [12] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In Advances in Neural Information Processing Systems, Vol. 33. Curran Associates, Inc., 9459–9474.
- [13] Zhen Lin, Shubhendu Trivedi, and Jimeng Sun. 2024. Generating with Confidence: Uncertainty Quantification for Black-box Large Language Models. Transactions on Machine Learning Research (2024).
- [14] Xinbei Ma, Yeyun Gong, Pengcheng He, Hai Zhao, and Nan Duan. 2023. Query Rewriting in Retrieval-Augmented Large Language Models. In Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics, Singapore, 5303–5315.
- [15] Timo Möller, Anthony Reina, Raghavan Jayakumar, and Malte Pietsch. 2020. COVID-QA: A Question Answering Dataset for COVID-19. In Proceedings of the 1st Workshop on NLP for COVID-19 at ACL 2020. Association for Computational Linguistics.
- [16] Noor Nashid, Mifta Sintaha, and Ali Mesbah. 2023. Retrieval-Based Prompt Selection for Code-Related Few-Shot Learning. In 2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE). 2450–2462.
- [17] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. Advances in neural information processing systems 35 (2022), 27730–27744.
- [18] Richard Yuanzhe Pang, Alicia Parrish, Nitish Joshi, Nikita Nangia, Jason Phang, Angelica Chen, Vishakh Padmakumar, Johnny Ma, Jana Thompson, He He, and Samuel Bowman. 2022. QuALITY: Question Answering with Long Input Texts, Yes!. In Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies. Association for Computational Linguistics, Seattle, United States, 5336–5358.
- [19] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. SQuAD: 100,000+ Questions for Machine Comprehension of Text. In Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics, Austin, Texas, 2383–2392.

- [20] Machel Reid, Nikolay Savinov, Denis Teplyashin, Dmitry Lepikhin, Timothy Lillicrap, Jean-baptiste Alayrac, Radu Soricut, Angeliki Lazaridou, Orhan Firat, Julian Schrittwieser, et al. 2024. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. https://arxiv.org/abs/2403.05530
- [21] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP). Association for Computational Linguistics, Hong Kong, China, 3982–3992.
- [22] Zhihong Shao, Yeyun Gong, Yelong Shen, Minlie Huang, Nan Duan, and Weizhu Chen. 2023. Enhancing Retrieval-Augmented Large Language Models with Iterative Retrieval-Generation Synergy. In Findings of the Association for Computational Linguistics: EMNLP 2023. Association for Computational Linguistics. Singapore, 9248–9274.
- [23] Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. 2019. CommonsenseQA: A Question Answering Challenge Targeting Commonsense Knowledge. In Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers). Association for Computational Linguistics, Minneapolis, Minnesota, 4149–4158.
- [25] Xidong Wang, Guiming Chen, Song Dingjie, Zhang Zhiyi, Zhihong Chen, Qingying Xiao, Junying Chen, Feng Jiang, Jianquan Li, Xiang Wan, Benyou Wang, and Haizhou Li. 2024. CMB: A Comprehensive Medical Benchmark in Chinese. In Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers). Mexico City, Mexico, 6184–6205.
- [26] Antoine Yang, Arsha Nagrani, Paul Hongsuck Seo, Antoine Miech, Jordi Pont-Tuset, Ivan Laptev, Josef Sivic, and Cordelia Schmid. 2023. Vid2Seq: Large-Scale Pretraining of a Visual Language Model for Dense Video Captioning. 2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) (2023), 10714–10726.
- [27] Xiao Yang, Kai Sun, Hao Xin, Yushi Sun, Nikita Bhalla, Xiangsen Chen, Sajal Choudhary, Rongze Daniel Gui, Ziran Will Jiang, Ziyu Jiang, Lingkun Kong, Brian Moran, Jiaqi Wang, Yifan Ethan Xu, An Yan, Chenyu Yang, Eting Yuan, Hanwen Zha, Nan Tang, Lei Chen, Nicolas Scheffer, Yue Liu, Nirav Shah, Rakesh Wanga, Anuj Kumar, Wen tau Yih, and Xin Luna Dong, 2024. CRAG Comprehensive RAG Benchmark. https://arxiv.org/abs/2406.04744
- [28] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. 2018. HotpotQA: A Dataset for Diverse, Explainable Multi-hop Question Answering. In Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics, Brussels, Belgium.
- [29] Wenhao Yu, Dan Iter, Shuohang Wang, Yichong Xu, Mingxuan Ju, Soumya Sanyal, Chenguang Zhu, Michael Zeng, and Meng Jiang. 2023. Generate rather than Retrieve: Large Language Models are Strong Context Generators. In The Eleventh International Conference on Learning Representations.
- [30] Jinming Zhao, Gholamreza Haffari, and Ehsan Shareghi. 2023. Generating Synthetic Speech from SpokenVocab for Speech Translation. In Findings of the Association for Computational Linguistics: EACL 2023. Association for Computational Linguistics, Dubrovnik, Croatia, 1975–1981.
- [31] Huaixiu Steven Zheng, Swaroop Mishra, Xinyun Chen, Heng-Tze Cheng, Ed H. Chi, Quoc V Le, and Denny Zhou. 2024. Take a Step Back: Evoking Reasoning via Abstraction in Large Language Models. In The Twelfth International Conference on Learning Representations.

# A Research Methods

# A.1 Prompt of Generating Answers with Domain in JSON format

You are an agent that only outputs JSON. You are given a Query and References. Do the following:

- 1. Determine the domain the query is about. The domain should be one of the following: "finance", "sports", "music", "movie", "encyclopedia". If none of the domains apply, use "other". Use "domain" as the key in the result json.
- Answer the question in as few words as possible. Please follow these guidelines when formulating your answer. If the question contains a false premise or assumption, answer

"invalid question". If you are uncertain or don't know the answer, respond with "I don't know". Use "answer" as the key in the result json.

```
### Question {query}
### References {references}
### Answer
```

# **Example Output**

```
{
    "domain": "movie", "answer": "Amy" }

# Step 1: Determine the domain
domain = "movie"

# Step 2: Answer the question
answer = "Amy"

# Create the result JSON
result = {"domain": domain, "answer": answer}
return result
```

Since sometimes the model generate additional text besides JSON output, additional code is written to extract the first JSON in the output to solve the problem.

# A.2 An example of hallucination

For example, a simple movie question "can you tell me the date that Randall Wallace was born?" (Ground truth: 1949-07-28), the model answers the same question repetitively:

- »According to the provided information, Randall Wallace was born on July 28, 1949:
- »According to my sources, Randall Wallace, the American screenwriter and director, was born on July 28, 1949;
- »I can confirm it now: Randall Wallace was born on November 7 1948
- »Randall Wallace's birthday is September 1, 1948.
- »I've got it nailed down now! Randall Wallace was indeed born on November 7, 1949
- »I finally have it exact: Randall Wallace was born on July 1, 1948 (no, wait!), I mean... September 1, 1949 (no), actually it is November 7, 1946. Yes, that's correct!